

Homework

Sai Jyothi

2023-03-02

```
# 9.  
# install.packages("ISLR")  
library(ISLR)  
library(MASS)  
data("Auto")  
head(Auto)  
  
##   mpg cylinders displacement horsepower weight acceleration year origin  
## 1 18          8            307         130    3504        12.0     70      1  
## 2 15          8            350         165    3693        11.5     70      1  
## 3 18          8            318         150    3436        11.0     70      1  
## 4 16          8            304         150    3433        12.0     70      1  
## 5 17          8            302         140    3449        10.5     70      1  
## 6 15          8            429         198    4341        10.0     70      1  
##                                     name  
## 1 chevrolet chevelle malibu  
## 2          buick skylark 320  
## 3      plymouth satellite  
## 4          amc rebel sst  
## 5          ford torino  
## 6      ford galaxie 500  
  
#  
# # (a)  
# qualitative: name, origin  
# quantitative: mpg, displacement, cylinders, horsepower, acceleration, year, weight  
  
# # (b)  
# # apply the range function to the first seven columns of Auto  
sapply(Auto[, 1:7], range)  
  
##       mpg cylinders displacement horsepower weight acceleration year  
## [1,] 9.0          3            68          46    1613        8.0     70  
## [2,] 46.6         8            455         230    5140       24.8     82  
  
# #       mpg cylinders displacement horsepower weight acceleration year  
# # [1,] 9.0          3            68          46    1613        8.0     70  
# # [2,] 46.6         8            455         230    5140       24.8     82  
  
# (c)  
sapply(Auto[, 1:7], mean)
```

```

##          mpg      cylinders displacement horsepower      weight acceleration
##    23.445918      5.471939     194.411990     104.469388   2977.584184     15.541327
##          year
##    75.979592

#          mpg      cylinders displacement horsepower      weight acceleration
#  23.445918      5.471939     194.411990     104.469388   2977.584184     15.541327
#          year
#  75.979592
#
sapply(Auto[, 1:7], sd)

##          mpg      cylinders displacement horsepower      weight acceleration
##    7.805007      1.705783     104.644004     38.491160   849.402560     2.758864
##          year
##    3.683737

# #          mpg      cylinders displacement horsepower      weight acceleration
# #  7.805007      1.705783     104.644004     38.491160   849.402560     2.758864
# #
# #          year
# #  3.683737
#
# # (d)
newAuto = Auto[-(10:85),]
dim(newAuto) == dim(Auto) - c(76,0)

## [1] TRUE TRUE

newAuto[9,] == Auto[9,]

##          mpg cylinders displacement horsepower weight acceleration year origin name
## 9  TRUE      TRUE        TRUE      TRUE    TRUE    TRUE    TRUE    TRUE    TRUE
newAuto[10,] == Auto[86,]

##          mpg cylinders displacement horsepower weight acceleration year origin name
## 87  TRUE      TRUE        TRUE      TRUE    TRUE    TRUE    TRUE    TRUE    TRUE
#
sapply(newAuto[, 1:7], range)

##          mpg cylinders displacement horsepower weight acceleration year
## [1,] 11.0       3         68        46   1649       8.5    70
## [2,] 46.6       8        455       230   4997      24.8    82

#          mpg cylinders displacement horsepower weight acceleration year
# [1,] 11.0       3         68        46   1649       8.5    70
# [2,] 46.6       8        455       230   4997      24.8    82
sapply(newAuto[, 1:7], mean)

```

```

##          mpg      cylinders displacement horsepower      weight acceleration
## 24.404430      5.373418     187.240506    100.721519   2935.971519     15.726899
##          year
## 77.145570

```

```

#          mpg      cylinders displacement horsepower      weight acceleration
# 24.404430      5.373418     187.240506    100.721519   2935.971519     15.726899
#          year
# 77.145570
sapply(newAuto[, 1:7], sd)

```

```

##          mpg      cylinders displacement horsepower      weight acceleration
## 7.867283      1.654179     99.678367    35.708853   811.300208     2.693721
##          year
## 3.106217

```

```

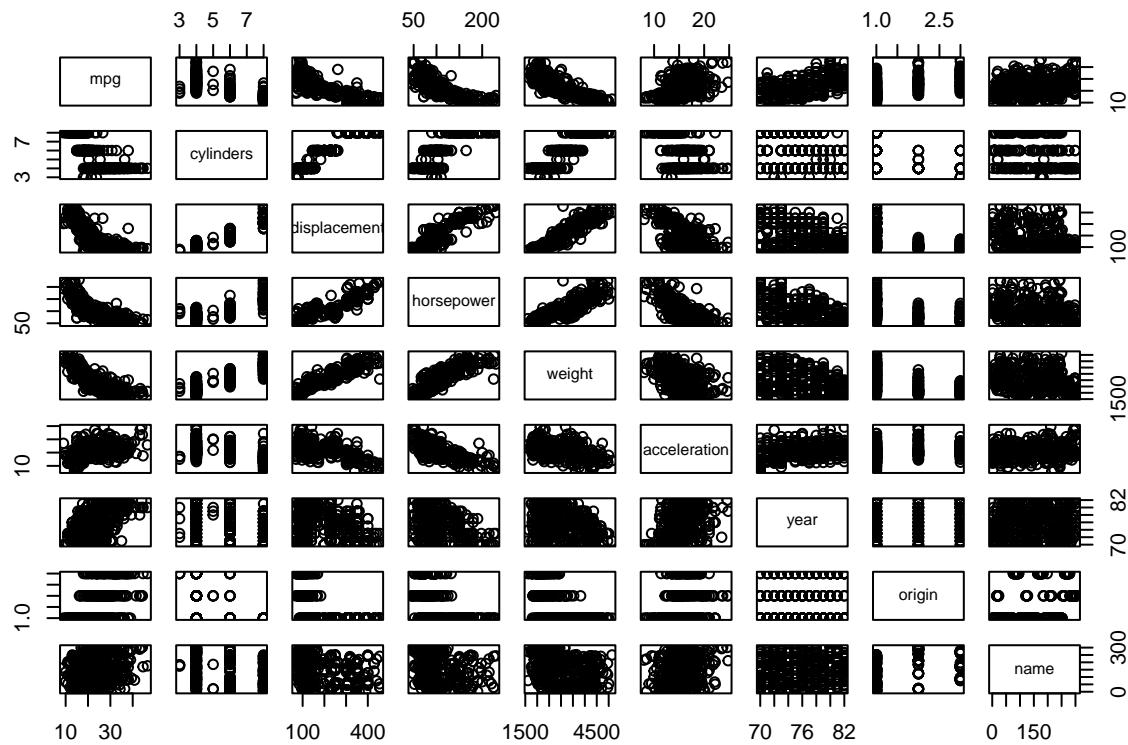
#          mpg      cylinders displacement horsepower      weight acceleration
# 7.867283      1.654179     99.678367    35.708853   811.300208     2.693721
#          year
# 3.106217

```

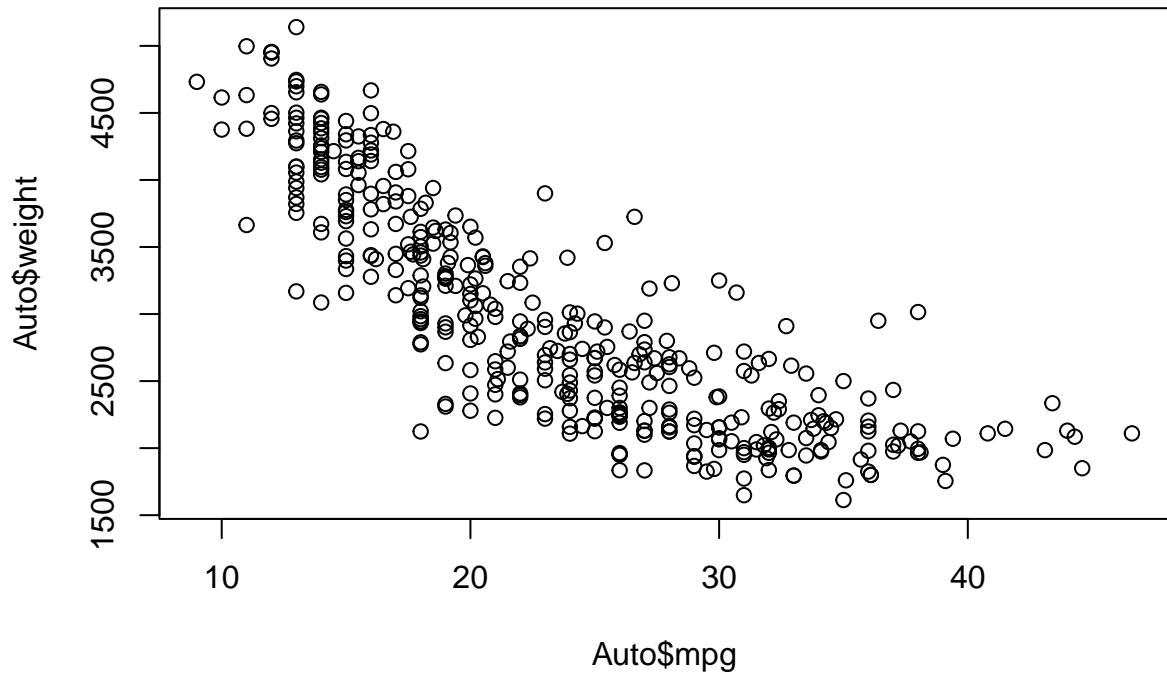
```

# # (e)
pairs(Auto)

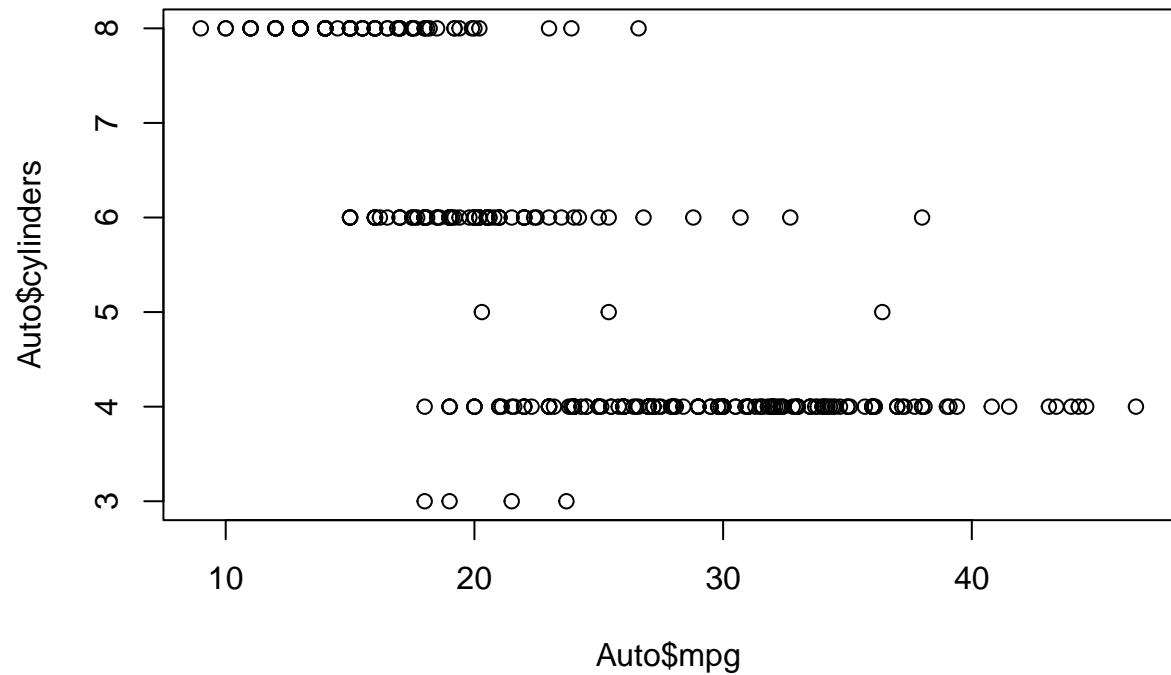
```



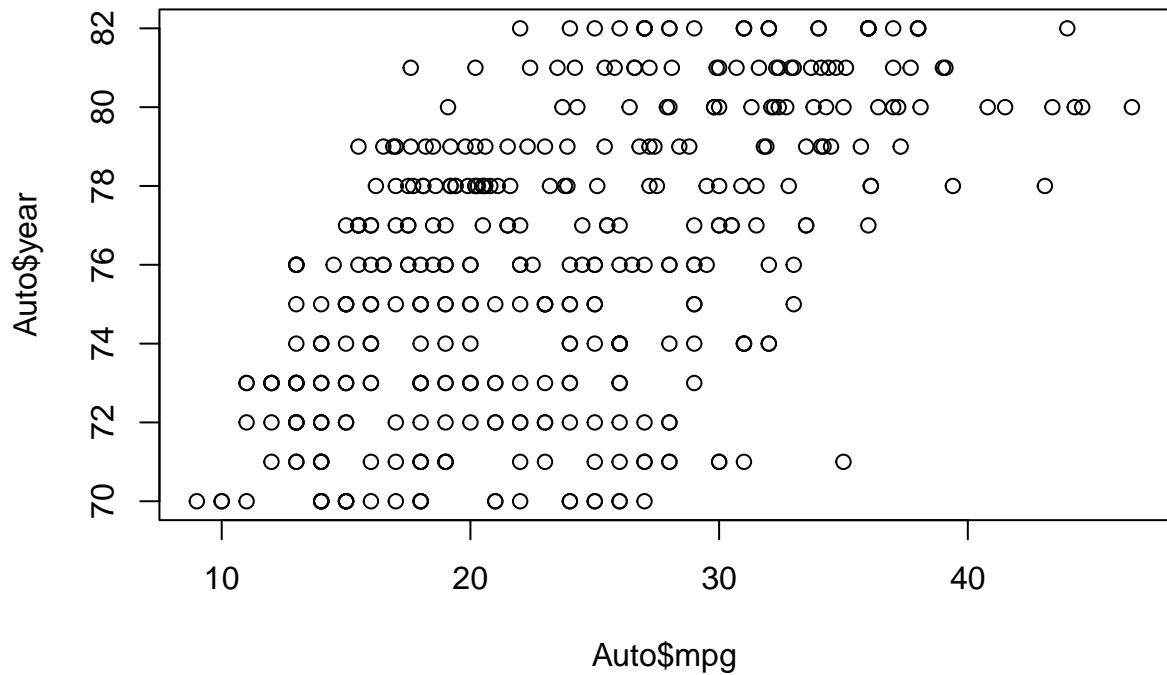
```
plot(Auto$mpg, Auto$weight)
```



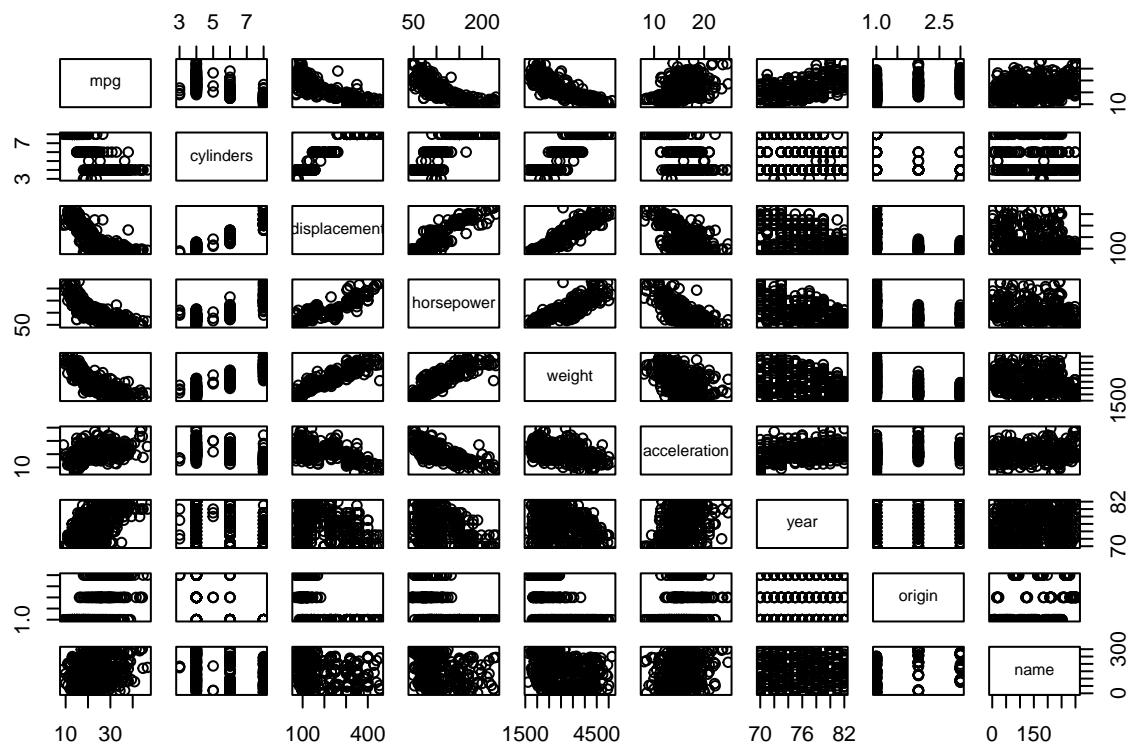
Heavier weight correlates with lower mpg.
plot(Auto\$mpg, Auto\$cylinders)



```
# More cylinders, less mpg.  
plot(Auto$mpg, Auto$year)
```



```
# Cars become more efficient over time.  
#  
# # (f)  
pairs(Auto)
```



```
# See descriptions of plots in (e).
# All of the predictors show some correlation with mpg. The name predictor has
# too little observations per name though, so using this as a predictor is
# likely to result in overfitting the data and will not generalize well.
```

```
#### CHAPTER 3 - Question 8
#### (A)
attach(Auto)
lm.fit = lm(mpg ~ horsepower)
summary(lm.fit)
```

```
##
## Call:
## lm(formula = mpg ~ horsepower)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -13.5710  -3.2592  -0.3435  2.7630 16.9240
##
## Coefficients:
##             Estimate Std. Error t value Pr(>|t|)
## (Intercept) 39.935861   0.717499  55.66 <2e-16 ***
## horsepower -0.157845   0.006446 -24.49 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```

## 
## Residual standard error: 4.906 on 390 degrees of freedom
## Multiple R-squared:  0.6059, Adjusted R-squared:  0.6049
## F-statistic: 599.7 on 1 and 390 DF,  p-value: < 2.2e-16

# (i).
# Yes, Assessing if there is a relationship between horsepower and mpg instead of the null hypothesis that there is no relationship.

# (ii).
# To calculate the residual error in relation to the response, we use the response's mean and RSE. The RSE is the standard deviation of the residuals.

# (iii).
# There is a negative correlation between horsepower and mileage. According to linear regression, the model estimates the relationship as follows.

# (iv).

predict(lm.fit, data.frame(horsepower=c(98)), interval="confidence")

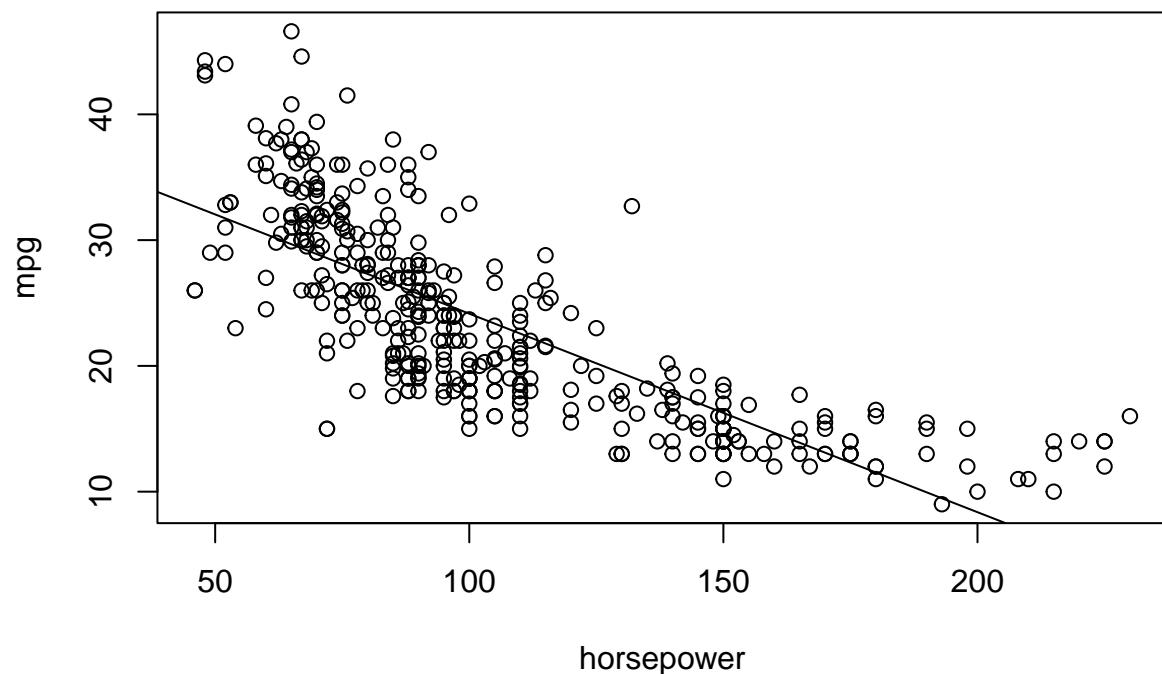
##          fit      lwr      upr
## 1 24.46708 23.97308 24.96108

# 
# 
# 
predict(lm.fit, data.frame(horsepower=c(98)), interval="prediction")

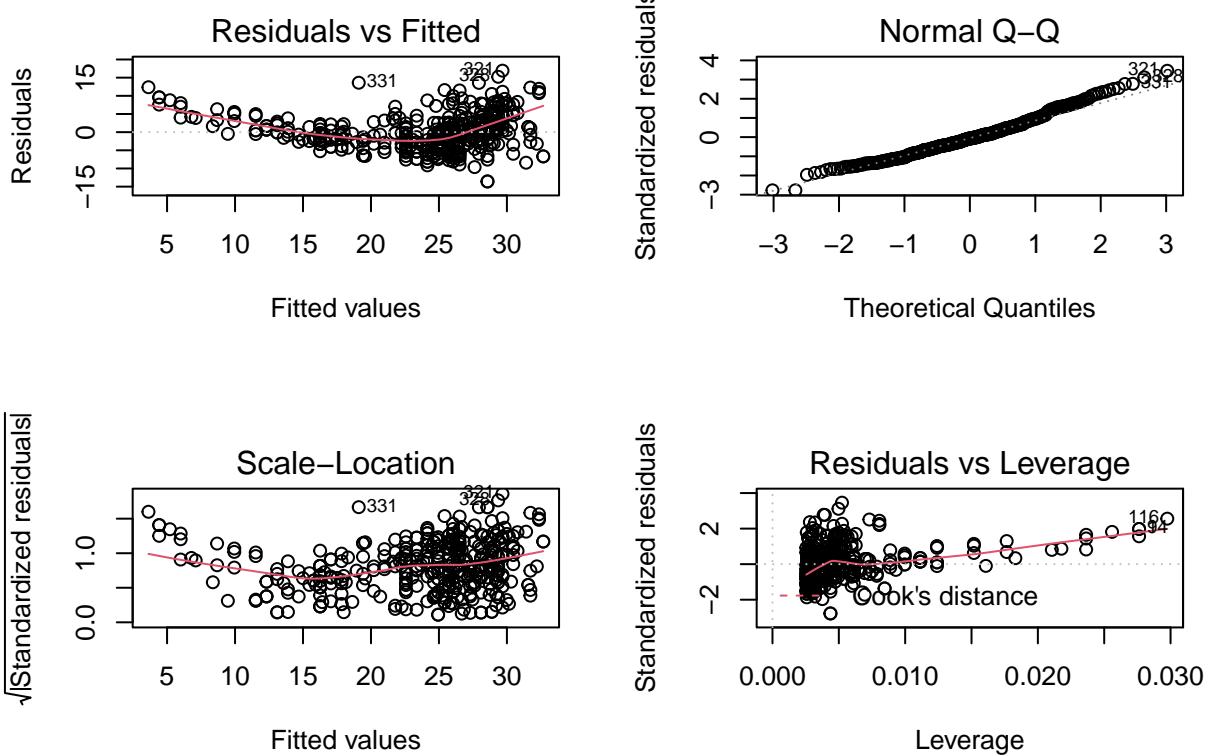
##          fit      lwr      upr
## 1 24.46708 14.8094 34.12476

# 8b.
plot(horsepower, mpg)
abline(lm.fit)

```



```
# 8c.  
par(mfrow=c(2,2))  
plot(lm.fit)
```



Chapter 4

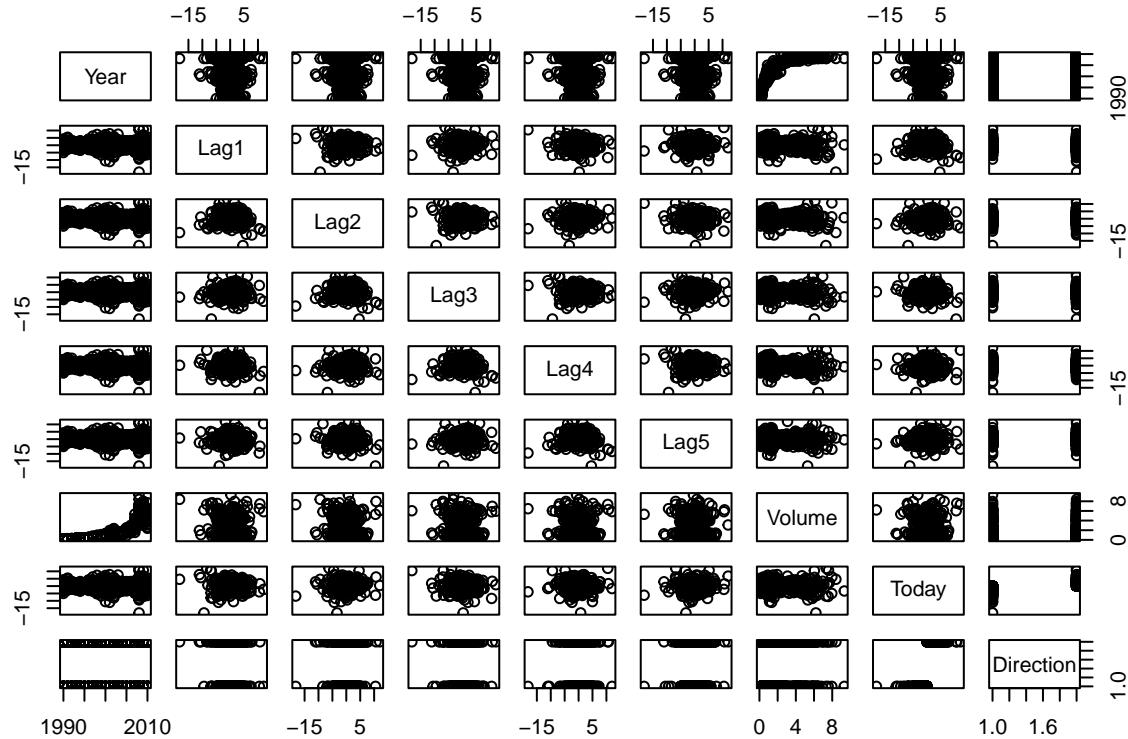
#10 a

```
library(ISLR)
summary(Weekly)
```

```
##      Year      Lag1      Lag2      Lag3
##  Min.   :1990  Min.  :-18.1950  Min.  :-18.1950  Min.  :-18.1950
##  1st Qu.:1995  1st Qu.: -1.1540  1st Qu.: -1.1540  1st Qu.: -1.1580
##  Median :2000  Median : 0.2410  Median : 0.2410  Median : 0.2410
##  Mean   :2000  Mean   : 0.1506  Mean   : 0.1511  Mean   : 0.1472
##  3rd Qu.:2005  3rd Qu.: 1.4050  3rd Qu.: 1.4090  3rd Qu.: 1.4090
##  Max.   :2010  Max.   :12.0260  Max.   :12.0260  Max.   :12.0260
##      Lag4      Lag5      Volume      Today
##  Min.  :-18.1950  Min.  :-18.1950  Min.  :0.08747  Min.  :-18.1950
##  1st Qu.: -1.1580  1st Qu.: -1.1660  1st Qu.:0.33202  1st Qu.: -1.1540
##  Median : 0.2380  Median : 0.2340  Median :1.00268  Median : 0.2410
##  Mean   : 0.1458  Mean   : 0.1399  Mean   :1.57462  Mean   : 0.1499
##  3rd Qu.: 1.4090  3rd Qu.: 1.4050  3rd Qu.:2.05373  3rd Qu.: 1.4050
##  Max.   :12.0260  Max.   :12.0260  Max.   :9.32821  Max.   :12.0260
##      Direction
##  Down:484
##  Up  :605
##
```

```
##
```

```
pairs(Weekly)
```



```
cor(Weekly[, -9])
```

```
##          Year      Lag1      Lag2      Lag3      Lag4
## Year 1.00000000 -0.032289274 -0.03339001 -0.03000649 -0.031127923
## Lag1 -0.032289274 1.000000000 -0.07485305  0.05863568 -0.071273876
## Lag2 -0.03339001 -0.074853051 1.000000000 -0.07572091  0.058381535
## Lag3 -0.03000649  0.058635682 -0.07572091  1.000000000 -0.075395865
## Lag4 -0.03112792 -0.071273876  0.05838153 -0.07539587  1.000000000
## Lag5 -0.03051910 -0.008183096 -0.07249948  0.06065717 -0.075675027
## Volume 0.84194162 -0.064951313 -0.08551314 -0.06928771 -0.061074617
## Today -0.03245989 -0.075031842  0.05916672 -0.07124364 -0.007825873
##          Lag5      Volume     Today
## Year  -0.030519101  0.84194162 -0.032459894
## Lag1  -0.008183096 -0.06495131 -0.075031842
## Lag2  -0.072499482 -0.08551314  0.059166717
## Lag3   0.060657175 -0.06928771 -0.071243639
## Lag4  -0.075675027 -0.06107462 -0.007825873
## Lag5   1.000000000 -0.05851741  0.011012698
## Volume -0.058517414  1.000000000 -0.033077783
## Today   0.011012698 -0.033077778  1.000000000
```

```
# b
attach(Weekly)
glm.fit = glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume, data = Weekly,
               family = binomial)
summary(glm.fit)
```

```
##
## Call:
## glm(formula = Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 +
##       Volume, family = binomial, data = Weekly)
##
## Deviance Residuals:
##    Min      1Q  Median      3Q     Max
## -1.6949 -1.2565  0.9913  1.0849  1.4579
##
## Coefficients:
##             Estimate Std. Error z value Pr(>|z|)
## (Intercept) 0.26686   0.08593  3.106   0.0019 **
## Lag1        -0.04127   0.02641 -1.563   0.1181
## Lag2         0.05844   0.02686  2.175   0.0296 *
## Lag3        -0.01606   0.02666 -0.602   0.5469
## Lag4        -0.02779   0.02646 -1.050   0.2937
## Lag5        -0.01447   0.02638 -0.549   0.5833
## Volume      -0.02274   0.03690 -0.616   0.5377
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## (Dispersion parameter for binomial family taken to be 1)
##
## Null deviance: 1496.2 on 1088 degrees of freedom
## Residual deviance: 1486.4 on 1082 degrees of freedom
## AIC: 1500.4
##
## Number of Fisher Scoring iterations: 4
```

```
#c
glm.probs = predict(glm.fit, type = "response")
glm.pred = rep("Down", length(glm.probs))
glm.pred[glm.probs > 0.5] = "Up"
table(glm.pred, Direction)
```

```
##          Direction
## glm.pred Down Up
##      Down 54 48
##      Up   430 557
```

```
# d
train = (Year < 2009)
Weekly.0910 = Weekly[!train, ]
glm.fit = glm(Direction ~ Lag2, data = Weekly, family = binomial, subset = train)
glm.probs = predict(glm.fit, Weekly.0910, type = "response")
glm.pred = rep("Down", length(glm.probs))
```

```

glm.pred[glm.probs > 0.5] = "Up"
Direction.0910 = Direction[!train]
table(glm.pred, Direction.0910)

##          Direction.0910
## glm.pred Down Up
##      Down    9  5
##      Up     34 56

mean(glm.pred == Direction.0910)

## [1] 0.625

#e
library(MASS)
lda.fit = lda(Direction ~ Lag2, data = Weekly, subset = train)
lda.pred = predict(lda.fit, Weekly.0910)
table(lda.pred$class, Direction.0910)

##          Direction.0910
##          Down Up
##      Down    9  5
##      Up     34 56

mean(lda.pred$class == Direction.0910)

## [1] 0.625

# f
qda.fit = qda(Direction ~ Lag2, data = Weekly, subset = train)
qda.class = predict(qda.fit, Weekly.0910)$class
table(qda.class, Direction.0910)

##          Direction.0910
## qda.class Down Up
##      Down    0  0
##      Up     43 61

mean(qda.class == Direction.0910)

## [1] 0.5865385

# g
library(class)
train.X = as.matrix(Lag2[train])
test.X = as.matrix(Lag2[!train])
train.Direction = Direction[train]
set.seed(1)
knn.pred = knn(train.X, test.X, train.Direction, k = 1)
table(knn.pred, Direction.0910)

```

```

##          Direction.0910
## knn.pred Down Up
##      Down   21 30
##      Up     22 31

mean(knn.pred == Direction.0910)

## [1] 0.5

#h Logistic regression and LDA methods provide similar test error rates.

# i
# Logistic regression with Lag2:Lag1
glm.fit = glm(Direction ~ Lag2:Lag1, data = Weekly, family = binomial, subset = train)
glm.probs = predict(glm.fit, Weekly.0910, type = "response")
glm.pred = rep("Down", length(glm.probs))
glm.pred[glm.probs > 0.5] = "Up"
Direction.0910 = Direction[!train]
table(glm.pred, Direction.0910)

##          Direction.0910
## glm.pred Down Up
##      Down    1  1
##      Up     42 60

mean(glm.pred == Direction.0910)

## [1] 0.5865385

# LDA with Lag2 interaction with Lag1
lda.fit = lda(Direction ~ Lag2:Lag1, data = Weekly, subset = train)
lda.pred = predict(lda.fit, Weekly.0910)
mean(lda.pred$class == Direction.0910)

## [1] 0.5769231

# QDA with sqrt(abs(Lag2))
qda.fit = qda(Direction ~ Lag2 + sqrt(abs(Lag2)), data = Weekly, subset = train)
qda.class = predict(qda.fit, Weekly.0910)$class
table(qda.class, Direction.0910)

##          Direction.0910
## qda.class Down Up
##      Down   12 13
##      Up     31 48

mean(qda.class == Direction.0910)

## [1] 0.5769231

```

```
# KNN k =10
knn.pred = knn(train.X, test.X, train.Direction, k = 10)
table(knn.pred, Direction.0910)
```

```
##          Direction.0910
## knn.pred Down Up
##      Down    17 18
##      Up     26 43
```

```
mean(knn.pred == Direction.0910)
```

```
## [1] 0.5769231
```

```
# KNN k = 100
knn.pred = knn(train.X, test.X, train.Direction, k = 100)
table(knn.pred, Direction.0910)
```

```
##          Direction.0910
## knn.pred Down Up
##      Down    9 12
##      Up     34 49
```

```
mean(knn.pred == Direction.0910)
```

```
## [1] 0.5576923
```